





CSE 322 : Computer Networks Sessional

NS3 PROJECT : TCP-AR (Adaptive Reno)

1705044



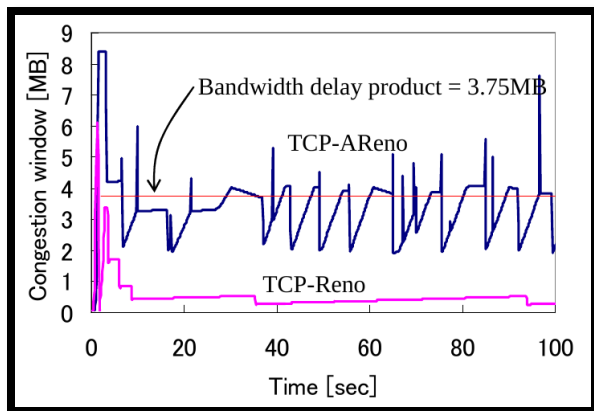
Paper Topology



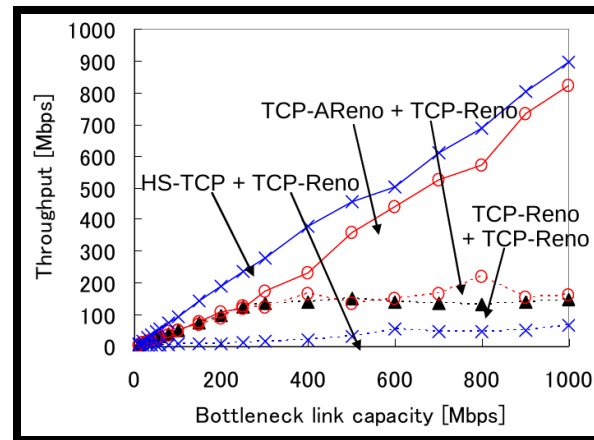
My Implementation

Paper Experiments

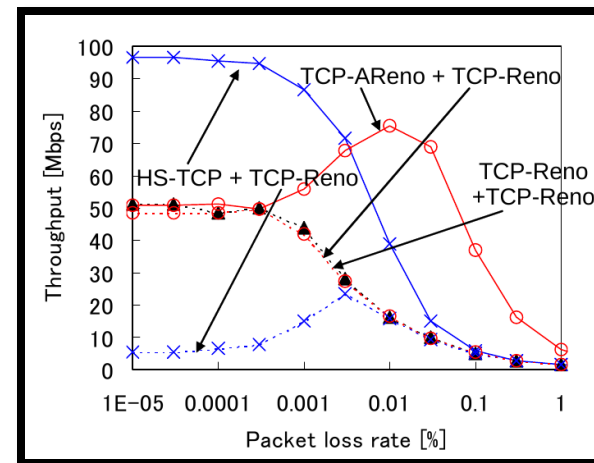
Throughput vs
bottleneck link
capacity



Congestion Window vs Time



Throughput vs
packet loss rate



Building Topology

```
// SETUP NODE AND DEVICE
// Create the point-to-point link helpers
PointToPointHelper bottleNeckLink;
bottleNeckLink.SetDeviceAttribute ("DataRate", StringValue (bottleNeckDataRate));
bottleNeckLink.SetChannelAttribute ("Delay", StringValue (bottleNeckDelay));

PointToPointHelper pointToPointLeaf;
pointToPointLeaf.SetDeviceAttribute ("DataRate", StringValue (senderDataRate));
pointToPointLeaf.SetChannelAttribute ("Delay", StringValue ("1ms"));

PointToPointDumbbellHelper d (nLeaf, pointToPointLeaf,
| | | | | | | | | | | | | | | | | | | | | |
                             nLeaf, pointToPointLeaf,
                             bottleNeckLink);
```

Building Topology

2 Different TCP variant for odd and even nodes

```
// INSTALL STACK
// tcp variant 1
Config::SetDefault ("ns3::TcpL4Protocol::SocketType", StringValue (tcpVariant1));
InternetStackHelper stack1;
for (uint32_t i = 0; i < d.LeftCount (); i+=2)
{
    stack1.Install (d.GetLeft (i)); // left leaves
}
for (uint32_t i = 0; i < d.RightCount (); i+=2)
{
    stack1.Install (d.GetRight (i)); // right leaves
}
stack1.Install (d.GetLeft ());
stack1.Install (d.GetRight ());

// tcp variant 2
Config::SetDefault ("ns3::TcpL4Protocol::SocketType", StringValue (tcpVariant2));
InternetStackHelper stack2;
for (uint32_t i = 1; i < d.LeftCount (); i+=2)
{
    stack2.Install (d.GetLeft (i)); // left leaves
}
for (uint32_t i = 1; i < d.RightCount (); i+=2)
{
    stack2.Install (d.GetRight (i)); // right leaves
}
```

Calculate Metrics

```
static void
CwndChange (Ptr<OutputStreamWrapper> stream, uint32_t oldCwnd, uint32_t newCwnd)
{
    // NS_LOG_UNCOND (Simulator::Now ().GetSeconds () << " " << newCwnd);
    *stream->GetStream () << Simulator::Now ().GetSeconds () << " " << newCwnd << std::endl;
}
```

```
std::ostringstream oss;
oss << output_folder << "/flow" << i+1 << ".cwnd";
AsciiTraceHelper asciiTraceHelper;
Ptr<OutputStreamWrapper> stream = asciiTraceHelper.CreateFileStream (oss.str());
ns3TcpSocket->TraceConnectWithoutContext ("CongestionWindow", MakeBoundCallback (&CwndChange, stream));
```

Calculate New Congestion Window for each flow

Calculate Metrics

```
Ptr<Ipv4FlowClassifier> classifier = DynamicCast<Ipv4FlowClassifier> (flowmon.GetClassifier ());
FlowMonitor::FlowStatsContainer stats = monitor->GetFlowStats ();

for (auto iter = stats.begin (); iter != stats.end (); ++iter) {
    Ipv4FlowClassifier::FiveTuple t = classifier->FindFlow (iter->first); |
    // classifier returns FiveTuple in correspondance to a flowID

    NS_LOG_UNCOND("----Flow ID:" <<iter->first);
    NS_LOG_UNCOND("Src Addr" <<t.sourceAddress << " -- Dst Addr "<< t.destinationAddress);
    NS_LOG_UNCOND("Sent Packets = " <<iter->second.txPackets);
    NS_LOG_UNCOND("Received Packets = " <<iter->second.rxPackets);
    NS_LOG_UNCOND("Lost Packets = " <<iter->second.txPackets-iter->second.rxPackets);
    NS_LOG_UNCOND("Packet delivery ratio = " <<iter->second.rxPackets*100.0/iter->second.txPackets << "%");
    NS_LOG_UNCOND("Packet loss ratio = " << (iter->second.txPackets-iter->second.rxPackets)*100.0/iter->second.txPackets << "%");
    NS_LOG_UNCOND(
        "Throughput = " <<iter->second.rxBytes * 8.0/(iter->second.timeLastRxPacket.GetSeconds()-iter->second.timeFirstTxPacket.GetSeconds())/1024<<"Kbps"
    );
}
```

Calculate Throughput and package drop rate for each flow

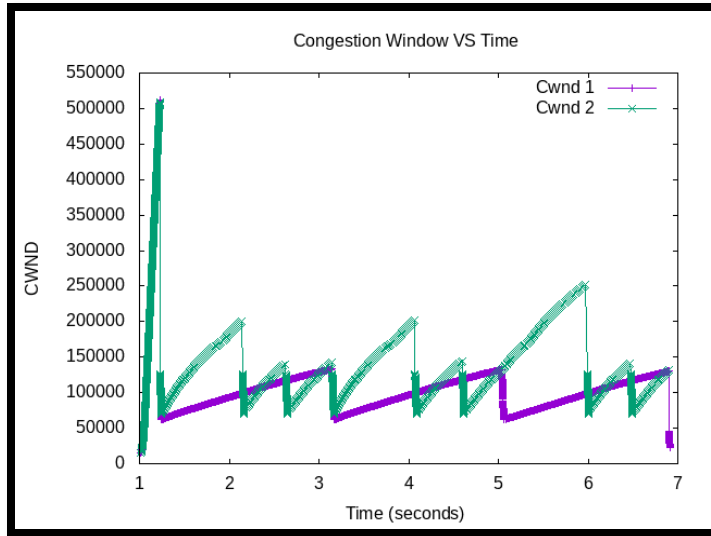
Full Pipeline

```
#!/bin/bash
rm -f -- temp/wired_dumbell/throughput1.txt
rm -f -- temp/wired_dumbell/throughput2.txt
touch temp/wired_dumbell/throughput1.txt
touch temp/wired_dumbell/throughput2.txt

for i in $(seq 10 10 90) #inclusive
do
    ./waf --run "scratch/wired_dumbell.cc --nLeaf=2 --bttlRate=${i}Mbps \
    --file1=temp/wired_dumbell/throughput1.txt --file2=temp/wired_dumbell/throughput2.txt"
    gnuplot -c temp/wired_dumbell/cwnd.plt "temp/wired_dumbell/cwnd${i}.png"
done

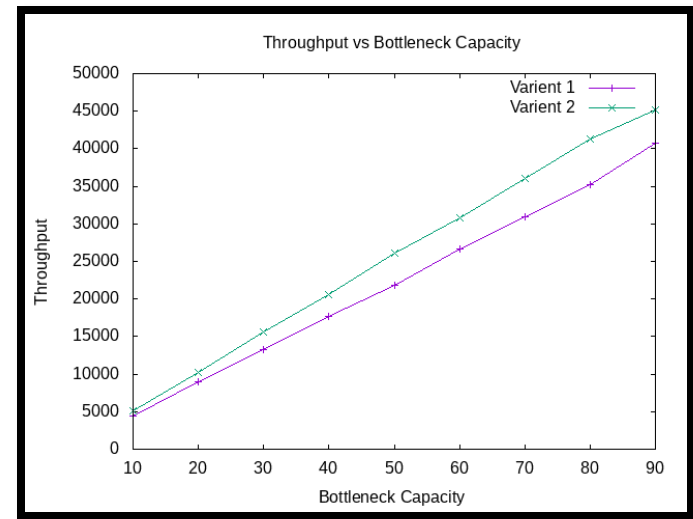
gnuplot temp/wired_dumbell/packetloss.plt
gnuplot temp/wired_dumbell/throughput.plt
```


Metric Graphs

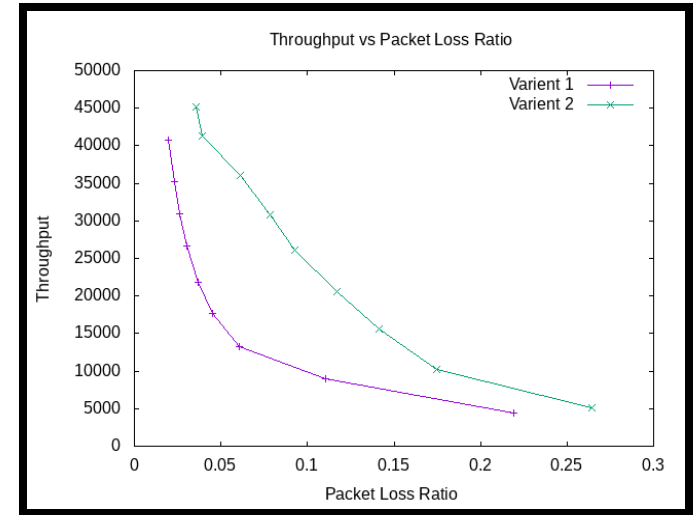


Congestion Window vs Time

Throughput vs bottleneck link capacity



Throughput vs packet loss rate






TCP-AdaptiveReno

Implementation Questions



Class Hierarchy



Variables

TCP Westwood

```
protected:
    TracedValue<double>    m_currentBW;           //!< Current value of the estimated BW
    double                 m_lastSampleBW;        //!< Last bandwidth sample
    double                 m_lastBW;              //!< Last bandwidth sample after being filtered
    enum ProtocolType      m_pType;              //!< 0 for Westwood, 1 for Westwood+
    enum FilterType        m_fType;              //!< 0 for none, 1 for Tustin

    uint32_t               m_ackedSegments;       //!< The number of segments ACKed between RTTs
    bool                   m_IsCount;             //!< Start keeping track of m_ackedSegments for Westwood+ if TRUE
    EventId                m_bwEstimateEvent;     //!< The BW estimation event for Westwood+
    Time                   m_lastAck;             //!< The last ACK time
```

TCP AdaptiveReno

```
protected:
    Time                   m_minRtt;              //!< Minimum RTT
    Time                   m_maxRtt;              //!< Maximum RTT (j th event)
    Time                   m_currentRtt;          //!< Current RTT
    Time                   m_prevMaxRtt;          //!< Previous Maximum RTT (j-1 th event)

    // Window calculations
    uint32_t               m_incWnd;              //!< Increment Window
    uint32_t               m_probeWnd;            //!< Probe Window
```

Functions

TCP NewReno

```
virtual std::string GetName () const;  
virtual uint32_t GetSsThresh (Ptr<const TcpSocketState> tcb, uint32_t bytesInFlight);  
virtual void IncreaseWindow (Ptr<TcpSocketState> tcb, uint32_t segmentsAacked);  
virtual void PktsAacked (Ptr<TcpSocketState> tcb, uint32_t segmentsAacked, const Time& rtt);  
virtual Ptr<TcpCongestionOps> Fork ();  
virtual void CwndEvent (Ptr<TcpSocketState> tcb, const TcpSocketState::TcpCaEvent_t event);
```

TCP Westwood

```
virtual uint32_t GetSsThresh (Ptr<const TcpSocketState> tcb, uint32_t bytesInFlight);  
virtual void PktsAacked (Ptr<TcpSocketState> tcb, uint32_t packetsAacked, const Time& rtt);  
void UpdateAackedSegments (int acked);  
void EstimateBW (const Time& rtt, Ptr<TcpSocketState> tcb);
```

TCP AdaptiveReno

```
double EstimateCongestionLevel(const Time& rtt, Ptr<TcpSocketState> tcb);  
void EstimateIncWnd(const Time& rtt, Ptr<TcpSocketState> tcb);
```

Modifications

```
/*
The function is called every time an ACK is received (only one time
also for cumulative ACKs) and contains timing information
*/
void
TcpAdaptiveReno::PktsAacked (Ptr<TcpSocketState> tcb, uint32_t packetsAacked,
                             const Time& rtt)
{
    NS_LOG_FUNCTION (this << tcb << packetsAacked << rtt);

    if (rtt.IsZero ())
    {
        NS_LOG_WARN ("RTT measured is zero!");
        return;
    }

    m_ackedSegments += packetsAacked;

    /*
    INITIALIZE AND SET VALUES FOR
    m_minRtt, m_maxRtt, m_currentRtt, m_prevMax
    m_incWnd, m_probeWnd
    */

    EstimateBW (rtt, tcb);
}
```

- Need to initialize and track several variables in ***PktsAacked()*** function

```
// initialize minRtt
if(m_minRtt.IsZero()) { m_minRtt = rtt; }
else if(rtt <= m_minRtt) { m_minRtt = rtt; }

// initialize maxRtt
if(m_maxRtt.IsZero()) { m_maxRtt = rtt; m_prevMaxRtt = m_maxRtt; }
else if(rtt >= m_maxRtt) { m_prevMaxRtt = m_maxRtt; m_maxRtt = rtt; }

// Update currentRTT
m_currentRtt = rtt;
```

Modifications

$$RTT_{cong}^j = (1-a)RTT_{cong}^{j-1} + aRTT^j$$

```
double
TcpAdaptiveReno::EstimateCongestionLevel(const Time &rtt, Ptr<TcpSocketState> tcb)
{
    /*
    float m_a = 0.75; // exponential smoothing factor
    double maxRtt = (1-m_a)*m_prevMaxRtt.GetSeconds() + m_a*m_maxRtt.GetSeconds();

    // NS_LOG_UNCOND("New maxRtt : "<<maxRtt);
    return std::min(
        (m_currentRtt.GetSeconds() - m_minRtt.GetSeconds()) / (maxRtt - m_minRtt.GetSeconds()),
        1.0
    );
    */
}
```

- Retrieve the current congestion level in ***EstimateCongestionFunction()*** function

$$c = \min\left(\frac{RTT - RTT_{\min}}{RTT_{cong} - RTT_{\min}}, 1\right)$$

Modifications

$$W_{inc}^{max} = B / M * MSS$$

$$W_{inc}(c) = W_{inc}^{max} / e^{\alpha c} + \beta c + \gamma$$

```
void
TcpAdaptiveReno::EstimateIncWnd(const Time& rtt, Ptr<TcpSocketState> tcb,
{
    /* double congestion = EstimateCongestionLevel();
    int scalingFactor_m = 10000;

    // m_currentBW; -> already calculated in packetsack?
    double m_maxIncWnd = m_currentBW / scalingFactor_m *
    | | | | static_cast<double> (tcb->m_segmentSize * tcb->m_segmentSize) ;

    double alpha = 2;
    double beta = 2 * m_maxIncWnd * ((1/alpha) - ((1/alpha + 1)/(std::exp(alpha)))));
    double gamma = 1 - (2 * m_maxIncWnd * ((1/alpha) - ((1/alpha + 0.5)/(std::exp(alpha)))));

    double temp = (m_maxIncWnd / std::exp(alpha * congestion)) + (beta * congestion) + gamma;
    m_incWnd = (int) temp;
}
```

- Retrieve the current increment window size in **EstimateIncWindow()** function

$$\beta = 2W_{inc}^{max} (1/\alpha - (1/\alpha + 1)/e^{\alpha})$$
$$\gamma = 1 - 2W_{inc}^{max} (1/\alpha - (1/\alpha + 1/2)/e^{\alpha}).$$

Modifications

$$\begin{aligned} W_{base} &= W_{base} + IMSS / W, \\ W_{probe} &= \max(W_{probe} + W_{inc} / W, 0) \end{aligned}$$

```

void
TcpAdaptiveReno::CongestionAvoidance (Ptr<TcpSocketState> tcb, uint32_t segmentsAcked)
{
    /* EstimateIncWnd(tcb);
    // base_window = USE NEW RENO IMPLEMENTATION
    double adder = static_cast<double> (tcb->m_segmentSize * tcb->m_segmentSize) /
    tcb->m_cWnd.Get ();
    adder = std::max (1.0, adder);
    m_baseWnd += static_cast<uint32_t> (adder);

    // change probe window
    // NS_LOG_UNCOND("incWnd "<m_incWnd<<" ; probe "<m_probeWnd);
    m_probeWnd = std::max(
        (double) (m_probeWnd + m_incWnd / (int)tcb->m_cWnd),
        (double) 0
    );
    tcb->m_cWnd = m_baseWnd + m_probeWnd;

```

- Calculate the current window size for the Congestion Avoidance Phase in **CongestionAvoidance()** function

Modifications

$$W_{base} = W * W_{dec} = W / (1+c), \quad W_{probe} =$$

```
uint32_t
TcpAdaptiveReno::GetSsThresh (Ptr<const TcpSocketState> tcb,
                             uint32_t bytesInFlight)
{
    /* double congestion = EstimateCongestionLevel();

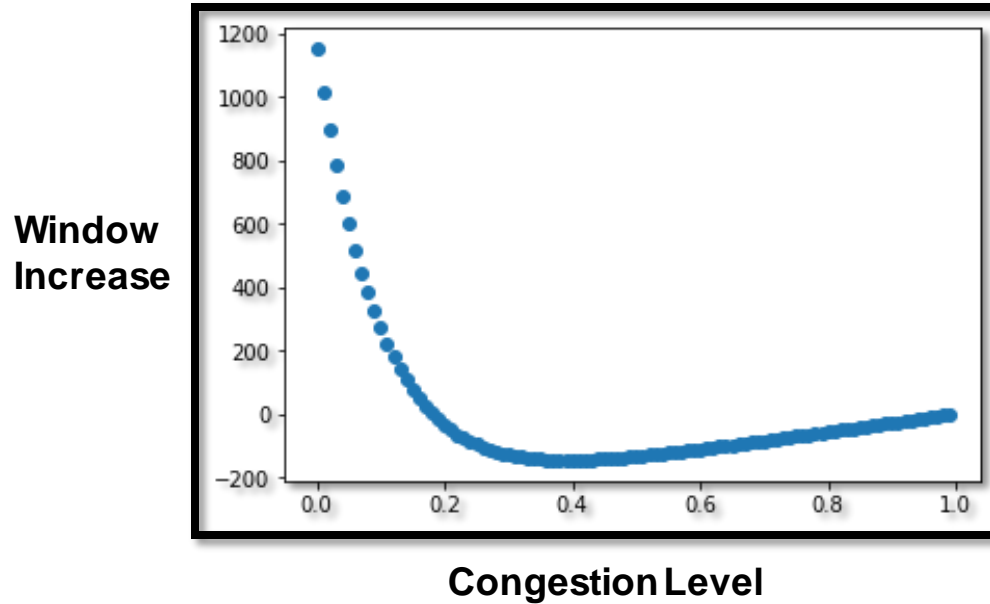
    uint32_t ssthresh = std::max (
        2*tcb->m_segmentSize,
        (uint32_t) (tcb->m_cWnd / (1.0+congestion))
    );

    m_baseWnd = ssthresh;
    m_probeWnd = 0;
    NS_LOG_LOGIC("new ssthresh : "<<ssthresh);

    return ssthresh;
```

- Calculate the Slow Start Threshold in the ***GetSsThresh()*** function

Query



- W_inc is never positive as congestion level is always > 0.4 in congestion avoidance phase

Reference

- H. Shimonishi and T. Murase, “**Improving efficiency friendliness tradeoffs of TCP congestion control algorithm,**” in Proc. IEEE GLOBECOM, 2005.
- Multiple NS3 Tutorials and documentations.

[Paper Link](#)



Thank you

Any Questions?

